Targeted Manipulation: Slope-Based Attacks on Financial Time Series Data

Dominik Luszczynski

Academic Supervisor: Professor Irene Huang

UNIVERSITY OF TORONTO

Introduction

Problem: The increased use of machine learning (ML) and artificial intelligence (AI) models in high-risk sectors, such as healthcare and finance, make it imperative that these predictive models are robust to adversarial attacks. Most contributions to adversarial attack research has been focused on image and text classification, while research about attacks on time series data, especially financial data, is still in its early stages [1].

Contribution: This research aims to build upon previous studies on adversarial attacks by introducing two slope-based targeted attacks on financial time series data, aimed to alter the temporal characteristics of a model's predictions. Furthermore, Generative Adversarial Networks (GANs) were experimented with to generate adversarial examples, based on a slope-based objective.

Robustness Criteria: The attack methods were tested on the novel N-HiTS architecture. In addition, the stealthiness of the attacks were measured against a 4-layered CNN was trained on varying sets of adversarial methods.

Method

Two novel slope attacks were implemented, the General Slope attack (GSA) and the Least Squares Slope Attack (LSSA), where each attack varies on how they compute the slope of the forecasted time series.

For hyperparameter scalars *c, d*, direction $t \in \{-1, 0, 1\}$, and a forecasted time series $(x_1, y_1), \ldots, (x_n, y_n)$, we compute the slopes *m,* and loss function as follows:

$$m_{GSA} = \frac{y_n - y_1}{x_n - x_1} \qquad m_{LSSA} = \frac{\sum_{i=0}^{N}(x - \bar{x})(y - \bar{y})}{(x - \bar{x})^2}$$

$$loss(m) = \begin{cases} ce^{-tdm}, & t \in \{-1, 1\} \\ cm^2, & t = 0 \end{cases}$$

An Adversarial GAN (A-GAN) was trained with stock data extracted similar to the N-HiTS model, specifically using the stock with the ticker *A*, where random continuous intervals of 99 days of log returns were selected for each sample of training data. A Conditional Wasserstein GAN was used for the A-GAN, conditioned with the corresponding 99 days of log returns. The condition is then concatenated with the noise vector in the feature dimension for both the generator and the critic.

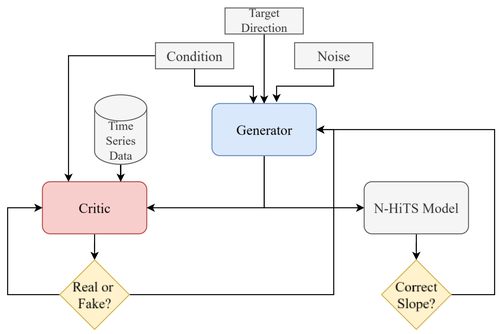

Results

GSA and LSSA Compared to Baseline Adversarial Attacks

Attack	*MAE*	*RMSE*	*MAPE*	*Gen. Slope*	*LS Slope*
Normal	2.15	2.72	3.82×10^{-2}	3.37×10^{-2}	2.22×10^{-2}
FGSM	2.57	3.21	4.51×10^{-2}	3.22×10^{-2}	2.34×10^{-2}
BIM	**3.38**	**3.99**	$\mathbf{5.68\times10^{-2}}$	3.48×10^{-2}	2.39×10^{-2}
MI-FGSM	3.37	**3.99**	5.67×10^{-2}	3.44×10^{-2}	2.39×10^{-2}
SIM	2.57	3.08	4.29×10^{-2}	3.37×10^{-2}	2.23×10^{-2}
TIM (Up)	2.49	3.21	4.52×10^{-2}	3.72×10^{-2}	2.00×10^{-2}
TIM (Down)	2.74	3.26	4.44×10^{-2}	3.32×10^{-2}	2.51×10^{-2}
GSA (Up)	2.26	2.88	4.03×10^{-2}	$\mathbf{6.76\times10^{-2}}$	2.77×10^{-2}
GSA (Down)	2.23	2.83	3.89×10^{-2}	$\mathbf{-1.68\times10^{-4}}$	1.75×10^{-2}
GSA (0)	2.30	2.93	4.01×10^{-2}	1.80×10^{-2}	2.00×10^{-2}
LSSA (Up)	2.49	3.10	4.26×10^{-2}	5.38×10^{-2}	$\mathbf{4.96\times10^{-2}}$
LSSA (Down)	2.71	3.33	4.63×10^{-2}	1.56×10^{-2}	$\mathbf{-5.04\times10^{-3}}$
LSSA (0)	2.68	3.31	4.55×10^{-2}	2.82×10^{-2}	1.29×10^{-2}

Average metrics for different attack methods performed on the first 300 days of each recording, with ε = 2% · median(adjprc). The best metrics are bolded.

Effect of GSA and LSSA on the N-HiTS Model

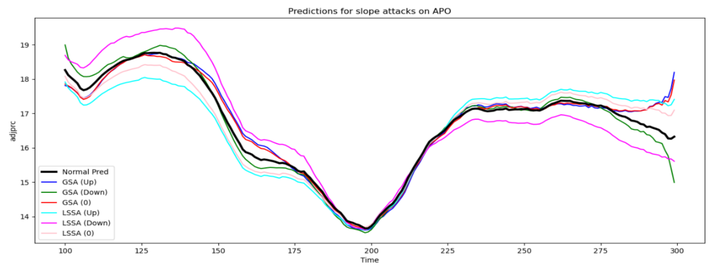

Example A-GAN Output and Effect on the N-HiTS Model

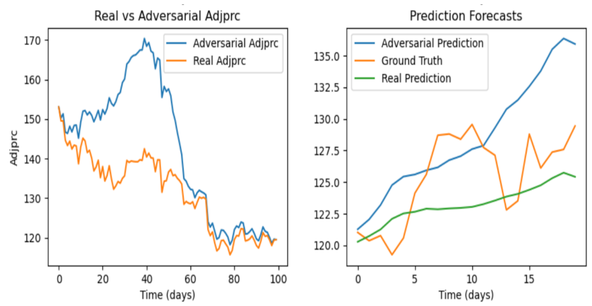

Stealthiness of GSA, LSSA and A-GAN

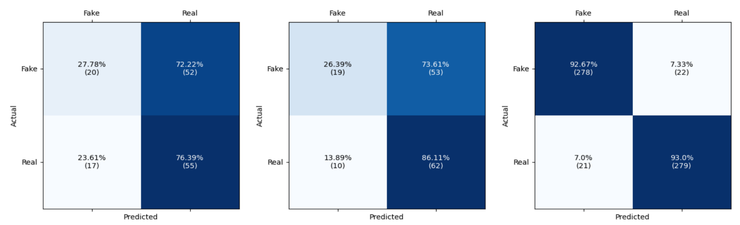

(a) GSA (b) LSSA (c) A-GAN

Confusion Matrices for a discriminator differentiating real from adversarial inputs. Both the GSA and LSSA confuse the discriminator, while A-GAN examples are easily identified as adversarial.

Key Findings and Conclusions

- Not only do the GSA and LSSA show a drastic effect on a state-of-the-art N-HiTS architecture trained for financial forecasting, they also have the ability to covertly bypass standard security measures.
- Incorporating the slope-based attack into the GAN architecture has shown that GANs are able to effectively generate synthetic adversarial examples.
- Although the A-GAN has shown to successfully increase the slope of the victim model's forecast, the A-GAN suffers from mode collapse, which occurs when the generator creates data with limited diversity [2].

Limitations and Future Work

- First, in order to solidify and prove the general effectiveness of the attack methods, several other ML models, like CNNs and LSTMs, should have been developed and used as the victim models, similar the study by Shen and Li [1].
- To prove the applicability of the slope-based attacks on time series data, these attacks should be implemented in other time series domains that use forecasting models, such as traffic and electricity usage.
- Given the A-GAN suffers from mode collapse, more experimentation should be performed to prevent it.

Further Information

For more information:

- GitHub: https://github.com/Dominik-luszcz/TargetedAdvGAN
- Email: dominik.luszczynski@mail.utoronto.ca

Acknowledgements

Thank you to Prof. Irene Huang, at the University of Toronto Scarborough, for supervising me during this CSCD94 project.

References

[1] Z. Shen and Y. Li, "Temporal characteristics-based adversarial attacks on time series forecasting," *Expert systems with applications*, vol. 264, Art. no. 125950, 2025, doi: 10.1016/j.eswa.2024.125950.

[2] Z. Dai, L. Zhao, K. Wang, and Y. Zhou, "Mode standardization: A practical countermeasure against mode collapse of GAN-based signal synthesis," *Applied soft computing*, vol. 150, Art. no. 111089, 2024, doi: 10.1016/j.asoc.2023.111089.